

# Part 1: Data Exploration

## Part 1.1: Understand the Raw Dataset

```
In [45]:
import pandas as pd
import numpy as np

churn_df = pd.read_csv('bank_data.csv')
churn_df.head()
```

Out[45]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.0
1	2	15647311	Hill	608	Spain	Female	41	1	83807.8
2	3	15619304	Onio	502	France	Female	42	8	159660.8
3	4	15701354	Boni	699	France	Female	39	1	0.0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.8

```
In [46]:
# check data info
churn_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   RowNumber       10000 non-null  int64
1   CustomerId      10000 non-null  int64
2   Surname         10000 non-null  object
3   CreditScore     10000 non-null  int64
4   Geography       10000 non-null  object
5   Gender         10000 non-null  object
6   Age            10000 non-null  int64
7   Tenure         10000 non-null  int64
8   Balance        10000 non-null  float64
9   NumOfProducts  10000 non-null  int64
10  HasCrCard       10000 non-null  int64
11  IsActiveMember  10000 non-null  int64
12  EstimatedSalary 10000 non-null  float64
13  Exited         10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [47]:
# check the unique value for each column
churn_df.nunique()
```

Out[47]:

RowNumber	10000
CustomerId	10000
Surname	2932
CreditScore	460
Geography	3
Gender	2
Age	70
Tenure	11
Balance	6382
NumOfProducts	4
HasCrCard	2
IsActiveMember	2
EstimatedSalary	9999
Exited	2

dtype: int64

## Part 1.2: Understand the Features

```
In [48]:
# check missing values
churn_df.isnull().sum()
```

Out[48]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

### Part 1.2.1: Understand the Numerical Features

In [49]:

```
churn_df[['CreditScore', 'Age', 'Tenure', 'NumOfProducts', 'Balance', 'EstimatedSalary']].describe()
```

Out[49]:

	CreditScore	Age	Tenure	NumOfProducts	Balance	EstimatedSal
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	1.530200	76485.889288	100090.239000
std	96.653299	10.487806	2.892174	0.581654	62397.405202	57510.492000
min	350.000000	18.000000	0.000000	1.000000	0.000000	11.580000
25%	584.000000	32.000000	3.000000	1.000000	0.000000	51002.110000
50%	652.000000	37.000000	5.000000	1.000000	97198.540000	100193.915000
75%	718.000000	44.000000	7.000000	2.000000	127644.240000	149388.247000
max	850.000000	92.000000	10.000000	4.000000	250898.090000	199992.480000

In [50]:

```
# check the features' distribution
# pandas.DataFrame.describe()
# boxplot, distplot, countplot
# the column 'Exited' is our target

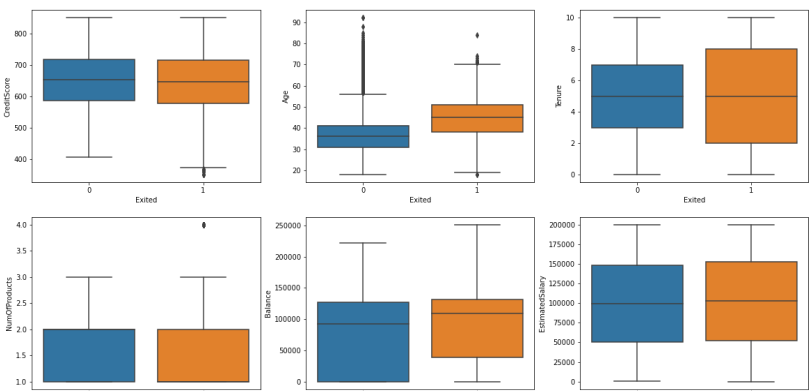
import matplotlib.pyplot as plt
import seaborn as sns
```

In [51]:

```
# boxplot for numerical feature
_,axss = plt.subplots(2,3, figsize=[20,10])
sns.boxplot(x='Exited', y='CreditScore', data=churn_df, ax=axss[0][0])
sns.boxplot(x='Exited', y='Age', data=churn_df, ax=axss[0][1])
sns.boxplot(x='Exited', y='Tenure', data=churn_df, ax=axss[0][2])
sns.boxplot(x='Exited', y='NumOfProducts', data=churn_df, ax=axss[1][0])
sns.boxplot(x='Exited', y='Balance', data=churn_df, ax=axss[1][1])
sns.boxplot(x='Exited', y='EstimatedSalary', data=churn_df, ax=axss[1][2])
```

Out[51]:

<AxesSubplot:xlabel='Exited', ylabel='EstimatedSalary'>



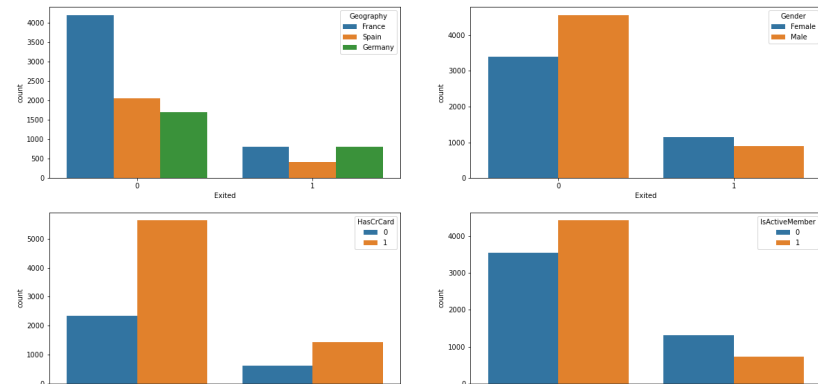
Part 1.2.2: Understand the Categorical Features

In [52]:

```
# countplot is a good choice here
_,axss = plt.subplots(2,2, figsize=[20,10])
sns.countplot(x='Exited', hue='Geography', data=churn_df, ax=axss[0][0])
sns.countplot(x='Exited', hue='Gender', data=churn_df, ax=axss[0][1])
sns.countplot(x='Exited', hue='HasCrCard', data=churn_df, ax=axss[1][0])
sns.countplot(x='Exited', hue='IsActiveMember', data=churn_df, ax=axss[1][1])
```

Out[52]:

<AxesSubplot:xlabel='Exited', ylabel='count'>



## Part 2: Feature Preprocessing

### Part 2.1: Extract Features

In [53]:

```
# Get feature space by dropping useless feature
# Obviously, these do not have a logical connection with churn
feature_drop = ['RowNumber', 'CustomerId', 'Surname', 'Exited']
X = churn_df.drop(feature_drop, axis=1)
X.head()
```

Out[53]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsAc
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

In [54]:

```
# There are two types of our features: categorical / numerical
X.dtypes
```

Out[54]:

```
CreditScore      int64
Geography        object
Gender           object
Age             int64
Tenure          int64
Balance         float64
NumOfProducts   int64
HasCrCard       int64
IsActiveMember  int64
EstimatedSalary float64
dtype: object
```

In [55]:

```
cat_cols = X.columns[X.dtypes == 'object']
num_cols = X.columns[(X.dtypes == 'float64') | (X.dtypes == 'int64')]
```

In [56]:

```
cat_cols
```

Out[56]:

```
Index(['Geography', 'Gender'], dtype='object')
```

In [57]:

```
num_cols
```

Out[57]:

```
Index(['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary'],
      dtype='object')
```

In [58]:

```
# Get target variable
y = churn_df['Exited']
```

In [59]:

```
# Split data into training(75%) and testing(25%) using model_selection function in sklearn

from sklearn import model_selection

X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0.25, stratify =
#stratified sampling

print('training data has ' + str(X_train.shape[0]) + ' observation with ' + str(X_train.shape[1]) + '
print('test data has ' + str(X_test.shape[0]) + ' observation with ' + str(X_test.shape[1]) + ' featu
# Show the size of both training of test data

training data has 7500 observation with 10 features
test data has 2500 observation with 10 features
```


In [60]:

```
# stratified sampling is used here to prevent extreme cases

X_train.head()
```

Out[60]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is
7971	633	Spain	Male	42	10	0.00	1	0	
9152	708	Germany	Female	23	4	71433.08	1	1	
6732	548	France	Female	37	9	0.00	2	0	
902	645	France	Female	48	7	90612.34	1	1	
2996	729	Spain	Female	45	7	91091.06	2	1	



Part 2.2: Encoding for Categorical Data

In [61]:

```
# One hot encoding
# Transform the categorical data of 'Geography' to numerical

from sklearn.preprocessing import OneHotEncoder

def OneHotEncoding(df, enc, categories):
    transformed = pd.DataFrame(enc.transform(df[categories]).toarray(), columns=enc.get_feature_names()
    return pd.concat([df.reset_index(drop=True), transformed], axis=1).drop(categories, axis=1)
# step1: Define a function to transform the result of OneHotCoding to dataframe
# step2: Substitute old column with new column

categories_1 = ['Geography']
enc_ohe = OneHotEncoder()
enc_ohe.fit(X_train[categories_1])
# Set up the One hot encoding size for 'Geography'


X_train = OneHotEncoding(X_train, enc_ohe, categories_1)
X_test = OneHotEncoding(X_test, enc_ohe, categories_1)
# Apply function to specific columns in both train and testing dataset
```

In [62]:

```
X_train.head()
```

Out[62]:

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	633	Male	42	10	0.00	1	0	1
1	708	Female	23	4	71433.08	1	1	0
2	548	Female	37	9	0.00	2	0	0
3	645	Female	48	7	90612.34	1	1	1
4	729	Female	45	7	91091.06	2	1	0



In [63]:

```
# Ordinal encoding
from sklearn.preprocessing import OrdinalEncoder

categories = ['Gender']
enc_oe = OrdinalEncoder()
enc_oe.fit(X_train[categories])


X_train[categories] = enc_oe.transform(X_train[categories])
X_test[categories] = enc_oe.transform(X_test[categories])
# Encoding 'Gender'
```

In [64]:

```
X_train.head()
```

Out[64]:

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	633	1.0	42	10	0.00	1	0	1
1	708	0.0	23	4	71433.08	1	1	0
2	548	0.0	37	9	0.00	2	0	0
3	645	0.0	48	7	90612.34	1	1	1
4	729	0.0	45	7	91091.06	2	1	0



Part 2.3: Standardize/Normalize Data

In [65]:

```
# Scale the data, using standardization
# Advantages:
# 1. speed up gradient descent
# 2. same scale
# 3. algorithm requirments


from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X_train[num_cols])
X_train[num_cols] = scaler.transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])
```

In [66]:

```
X_train.head()
```

Out[66]:

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveM
0	-0.172985	1.0	0.289202	1.731199	-1.218916	-0.912769	-1.542199	0.5
1	0.602407	0.0	-1.509319	-0.341156	-0.076977	-0.912769	0.648425	-1.0
2	-1.051762	0.0	-0.184093	1.385806	-1.218916	0.796109	-1.542199	-1.0
3	-0.048922	0.0	0.857156	0.695022	0.229625	-0.912769	0.648425	0.5
4	0.819517	0.0	0.573179	0.695022	0.237278	0.796109	0.648425	-1.0



Part 3: Model Training and Result Evaluation

Part 3.1: Model Training

In [67]:

```
# Build models
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression

# Logistic Regression
classifier_logistic = LogisticRegression()

# K Nearest Neighbors
classifier_KNN = KNeighborsClassifier()

# Random Forest
classifier_RF = RandomForestClassifier()
```

In [92]:

```
# Train the model
classifier_logistic.fit(X_train, y_train)
```

Out[92]:

LogisticRegression()

In [93]:

```
# Prediction of test data
classifier_logistic.predict(X_test)
```

Out[93]:

array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

In [94]:

```
# Accuracy of test data
classifier_logistic.score(X_test, y_test)
```

Out[94]:

0.8088

In [71]:

```
# Use 5-fold Cross Validation to get the accuracy for different models
model_names = ['Logistic Regression', 'KNN', 'Random Forest']
model_list = [classifier_logistic, classifier_KNN, classifier_RF]
count = 0

for classifier in model_list:
    cv_score = model_selection.cross_val_score(classifier, X_train, y_train, cv=5)
    print(cv_score)
    print('Model accuracy of ' + model_names[count] + ' is ' + str(cv_score.mean()))
    count += 1
```

```
[0. 81933333 0. 80666667 0. 80666667 0. 80933333 0. 82      ]
Model accuracy of Logistic Regression is 0.8124
[0. 84133333 0. 84066667 0. 83      0. 83066667 0. 84      ]
Model accuracy of KNN is 0.8365333333333334
[0. 87666667 0. 86266667 0. 85466667 0. 85866667 0. 86266667]
Model accuracy of Random Forest is 0.8630666666666666
```

## Part 3.2: Use Grid Search to Find Optimal Hyperparameters

In [72]:

```
# Loss/cost function -->  $(wx + b - y)^2 + \lambda * |w|$  -->  $\lambda$  is a hyperparameter
```

In [95]:

```
from sklearn.model_selection import GridSearchCV

# helper function for printing out grid search results
def print_grid_search_metrics(gs):
    print("Best score: " + str(gs.best_score_))
    print("Best parameters set:")
    best_parameters = gs.best_params_
    for param_name in sorted(best_parameters.keys()):
        print(param_name + ':' + str(best_parameters[param_name]))
```

### Part 3.2.1: Find Optimal Hyperparameters - LogisticRegression-lambda

In [96]:

```
# Possible hyperparamter options for Logistic Regression Regularization
# Penalty is choosed from L1 or L2
# C is the 1/lambda value(weight) for L1 and L2
# solver: algorithm to find the weights that minimize the cost function

# ('l1', 0.01) ('l1', 0.05) ('l1', 0.1) ('l1', 0.2) ('l1', 1)
# ('l2', 0.01) ('l2', 0.05) ('l2', 0.1) ('l2', 0.2) ('l2', 1)
parameters = {
    'penalty': ('l1', 'l2'),
    'C': (0.01, 0.05, 0.1, 0.2, 1)
}
Grid_LR = GridSearchCV(LogisticRegression(solver='liblinear'), parameters, cv=5)
Grid_LR.fit(X_train, y_train)
```

Out[96]:

```
GridSearchCV(cv=5, estimator=LogisticRegression(solver='liblinear'),
             param_grid={'C': (0.01, 0.05, 0.1, 0.2, 1),
                          'penalty': ('l1', 'l2')})
```

In [105]:

```
# the best hyperparameter combination
# C = 1/lambda
print_grid_search_metrics(Grid_LR)
```

```
Best score: 0.8125333333333333
Best parameters set:
C:1
penalty:l1
```

In [106]:

```
# best model
best_LR_model = Grid_LR.best_estimator_
best_LR_model
```

Out[106]:

```
LogisticRegression(C=1, penalty='l1', solver='liblinear')
```

In [107]:

```
best_LR_model.predict(X_test)
```

Out[107]:

```
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

In [108]:


```
best_LR_model.score(X_test, y_test)
```

Out[108]:

```
0.8092
```

In [109]:

```
LR_models = pd.DataFrame(Grid_LR.cv_results_)
res = (LR_models.pivot(index='param_penalty', columns='param_C', values='mean_test_score'))
_ = sns.heatmap(res, cmap='magma')
```



### Part 3.2.2: Find Optimal Hyperparameters: KNN-C

In [110]:

```
# Possible hyperparamter options for KNN
# Choose k
parameters = {
    'n_neighbors': [1, 3, 5, 7, 9]
}
Grid_KNN = GridSearchCV(KNeighborsClassifier(), parameters, cv=5)
Grid_KNN.fit(X_train, y_train)
```

Out[110]:

```
GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
             param_grid={'n_neighbors': [1, 3, 5, 7, 9]})
```

In [111]:

```
# best k
print_grid_search_metrics(Grid_KNN)
```

Best score: 0.8433333333333334  
Best parameters set:  
n\_neighbors:9

In [112]:

```
best_KNN_model = Grid_KNN.best_estimator_
```

### Part 3.2.3: Find Optimal Hyperparameters: Random Forest-depth

In [117]:

```
# Possible hyperparamter options for Random Forest
# Choose the number of trees
parameters = {
    'n_estimators': [60, 80, 100],
    'max_depth': [1, 5, 10]
}
Grid_RF = GridSearchCV(RandomForestClassifier(), parameters, cv=5)
Grid_RF.fit(X_train, y_train)
```

Out[117]:

```
GridSearchCV(cv=5, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 5, 10],
                         'n_estimators': [60, 80, 100]})
```

In [114]:

```
# best number of tress
print_grid_search_metrics(Grid_RF)
```

Best score: 0.8664000000000002  
Best parameters set:  
max\_depth:10  
n\_estimators:80

In [115]:

```
# best random forest
best_RF_model = Grid_RF.best_estimator_
best_RF_model
```

Out[115]:

```
RandomForestClassifier(max_depth=10, n_estimators=80)
```

## Part 3.3: Model Evaluation - Confusion Matrix (Precision, Recall, Accuracy)

TP: correctly labeled real churn

Precision (PPV, positive predictive value):  $tp / (tp + fp)$ ; Total number of true predictive churn divided by the total number of predictive churn; High Precision means low fp, not many return users were predicted as churn users.

Recall (sensitivity, hit rate, true positive rate):  $tp / (tp + fn)$  Predict most positive or churn user correctly. High recall means low fn, not many churn users were predicted as return users.

In [118]:

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

# calculate accuracy, precision and recall, [[tn, fp], []]
def cal_evaluation(classifier, cm):
    tn = cm[0][0]
    fp = cm[0][1]
    fn = cm[1][0]
    tp = cm[1][1]
    accuracy = (tp + tn) / (tp + fp + fn + tn + 0.0)
    precision = tp / (tp + fp + 0.0)
    recall = tp / (tp + fn + 0.0)
    print (classifier)
    print ("Accuracy is: " + str(accuracy))
    print ("precision is: " + str(precision))
    print ("recall is: " + str(recall))
    print ()

# print out confusion matrices
def draw_confusion_matrices(confusion_matrices):
    class_names = ['Not', 'Churn']
    for cm in confusion_matrices:
        classifier, cm = cm[0], cm[1]
        cal_evaluation(classifier, cm)
```

In [119]:

```
# Confusion matrix, accuracy, precison and recall for random forest and logistic regression
confusion_matrices = [
    ("Random Forest", confusion_matrix(y_test, best_RF_model.predict(X_test))),
    ("Logistic Regression", confusion_matrix(y_test, best_LR_model.predict(X_test))),
    ("K nearest neighbor", confusion_matrix(y_test, best_KNN_model.predict(X_test)))
]

draw_confusion_matrices(confusion_matrices)
```

Random Forest  
Accuracy is: 0.8604  
precision is: 0.8076923076923077  
recall is: 0.412573673870334

Logistic Regression  
Accuracy is: 0.8092  
precision is: 0.5963855421686747  
recall is: 0.1944990176817289

K nearest neighbor  
Accuracy is: 0.8428  
precision is: 0.7283464566929134  
recall is: 0.36345776031434185

## Part 3.4: Model Evaluation - ROC & AUC

### Part 3.4.1: ROC of RF Model

In [120]:

```
from sklearn.metrics import roc_curve
from sklearn import metrics

# Use predict_proba to get the probability results of Random Forest
y_pred_rf = best_RF_model.predict_proba(X_test)[: , 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)
```

In [121]:


```
best_RF_model.predict_proba(X_test)
```

Out[121]:

```
array([[0.75138239, 0.24861761],
       [0.92837177, 0.07162823],
       [0.73237493, 0.26762507],
       ...,
       [0.84666041, 0.15333959],
       [0.92533831, 0.07466169],
       [0.90282888, 0.09717112]])
```

In [122]:

```
# ROC curve of Random Forest result
import matplotlib.pyplot as plt
plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rf, tpr_rf, label='RF')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - RF model')
plt.legend(loc='best')
plt.show()
```





In [123]:

```
from sklearn import metrics

# AUC score
# The probability that a randomly-chosen positive example is ranked more highly than a randomly-chosen negative example
metrics.auc(fpr_rf, tpr_rf)
```

Out[123]:

0.8459640089637159

## Part 3.4.2: ROC of Logistic Regression Model

In [124]:

```
# Use predict_proba to get the probability results of Logistic Regression
y_pred_lr = best_LR_model.predict_proba(X_test)[: , 1]
fpr_lr, tpr_lr, thresh = roc_curve(y_test, y_pred_lr)
```

In [125]:


```
best_LR_model.predict_proba(X_test)
```

Out[125]:

```
array([[0.82435829, 0.17564171],
       [0.9317178 , 0.0682822 ],
       [0.85520934, 0.14479066],
       ...,
       [0.71449535, 0.28550465],
       [0.89278331, 0.10721669],
       [0.85561097, 0.14438903]])
```

In [126]:

```
# ROC Curve
plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_lr, tpr_lr, label='LR')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - LR Model')
plt.legend(loc='best')
plt.show()
```



In [127]:

```
# AUC score
metrics.auc(fpr_lr, tpr_lr)
```

Out[127]:

0.7722008369687169

## Part 3.4.3: ROC of KNN Model

In [128]:

```
# Use predict_proba to get the probability results of Logistic Regression
y_pred_lr = best_KNN_model.predict_proba(X_test)[: , 1]
fpr_lr, tpr_lr, thresh = roc_curve(y_test, y_pred_lr)
```

In [129]:


```
best_KNN_model.predict_proba(X_test)
```

Out[129]:

```
array([[1.          , 0.          ],
       [1.          , 0.          ],
       [0.88888889, 0.11111111],
       ...,
       [0.77777778, 0.22222222],
       [1.          , 0.          ],
       [1.          , 0.          ]])
```

In [130]:

```
# ROC Curve
plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_lr, tpr_lr, label='KNN')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - KNN Model')
plt.legend(loc='best')
plt.show()
```



In [131]:

```
# AUC score
metrics.auc(fpr_lr, tpr_lr)
```

Out[131]:

0.7986385690420251

## Part 4: Model Extra Functionality

### Part 4.1: Logistic Regression Model

In [132]:

```
X_with_corr = X.copy()
```

```
X_with_corr = OneHotEncoding(X_with_corr, enc_ohe, ['Geography'])
X_with_corr['Gender'] = enc_oe.transform(X_with_corr[['Gender']])
X_with_corr.head()
```

Out[132]:

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	0.0	42	2	0.00	1	1	1
1	608	0.0	41	1	83807.86	1	0	1
2	502	0.0	42	8	159660.80	3	1	0
3	699	0.0	39	1	0.00	2	0	0
4	850	0.0	43	2	125510.82	1	1	1

In [133]:

```
# Add L1 regularization to logistic regression
# Check the coef for feature selection
```

```
scaler = StandardScaler()
X_l1 = scaler.fit_transform(X_with_corr)
LRmodel_l1 = LogisticRegression(penalty="l1", C = 0.04, solver='liblinear')
LRmodel_l1.fit(X_l1, y)
```

```
indices = np.argsort(abs(LRmodel_l1.coef_[0]))[::-1]
```

```
print ("Logistic Regression (L1) Coefficients")
for ind in range(X_with_corr.shape[1]):
    print ("{0} : {1}".format(X_with_corr.columns[indices[ind]], round(LRmodel_l1.coef_[0][indices[ind]]
```

Logistic Regression (L1) Coefficients

Age : 0.7307

IsActiveMember : -0.5046

Geography\_Germany : 0.3121

Gender : -0.2409

Balance : 0.1509

CreditScore : -0.0457

NumOfProducts : -0.0439

Tenure : -0.0271

EstimatedSalary : 0.0092

Geography\_France : -0.0042

HasCrCard : -0.0022

Geography\_Spain : 0.0

In [134]:

```
# Add L2 regularization to logistic regression
# Check the coef for feature selection

np.random.seed()
scaler = StandardScaler()
X_l2 = scaler.fit_transform(X_with_corr)
LRmodel_l2 = LogisticRegression(penalty="l2", C = 0.1, solver='liblinear', random_state=42)
LRmodel_l2.fit(X_l2, y)
LRmodel_l2.coef_[0]

indices = np.argsort(abs(LRmodel_l2.coef_[0]))[::-1]

print ("Logistic Regression (L2) Coefficients")
for ind in range(X_with_corr.shape[1]):
    print ("{0} : {1}".format(X_with_corr.columns[indices[ind]], round(LRmodel_l2.coef_[0][indices[ind]], 4)))
```

Logistic Regression (L2) Coefficients  
Age : 0.751  
IsActiveMember : -0.5272  
Gender : -0.2591  
Geography\_Germany : 0.2279  
Balance : 0.162  
Geography\_France : -0.1207  
Geography\_Spain : -0.089  
CreditScore : -0.0637  
NumOfProducts : -0.0586  
Tenure : -0.0452  
EstimatedSalary : 0.0272  
HasCrCard : -0.0199

## Part 4.2: Random Forest Model - Feature Importance Discussion

In [135]:

```
X_RF = X.copy()

X_RF = OneHotEncoding(X_RF, enc_oe, ['Geography'])
X_RF['Gender'] = enc_oe.transform(X_RF[['Gender']])

X_RF.head()
```

Out[135]:

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	0.0	42	2	0.00	1	1	1
1	608	0.0	41	1	83807.86	1	0	1
2	502	0.0	42	8	159660.80	3	1	0
3	699	0.0	39	1	0.00	2	0	0
4	850	0.0	43	2	125510.82	1	1	1

In [136]:

```
# Check feature importance of random forest for feature selection
forest = RandomForestClassifier()
forest.fit(X_RF, y)

importances = forest.feature_importances_

indices = np.argsort(importances)[::-1]

# Print the feature ranking
print("Feature importance ranking by Random Forest Model:")
for ind in range(X.shape[1]):
    print ("{0} : {1}".format(X_RF.columns[indices[ind]], round(importances[indices[ind]], 4)))
```

Feature importance ranking by Random Forest Model:  
Age : 0.2395  
EstimatedSalary : 0.1469  
CreditScore : 0.1449  
Balance : 0.1403  
NumOfProducts : 0.1271  
Tenure : 0.0827  
IsActiveMember : 0.0423  
Geography\_Germany : 0.0206  
Gender : 0.0185  
HasCrCard : 0.0181